

ChE352
Numerical Techniques for Chemical Engineers
Professor Stevenson

Lecture 3

Numpy arrays

```
import numpy as np
x = np.array([1, 4, 3])
y = np.array([[1, 4, 3], [9, 2, 7]])
print('x.shape:', x.shape, ' x.size:', x.size)
print('y.shape:', y.shape, ' y.size:', y.size)
```

x.shape: (3,) x.size: 3

y.shape: (2, 3) y.size: 6

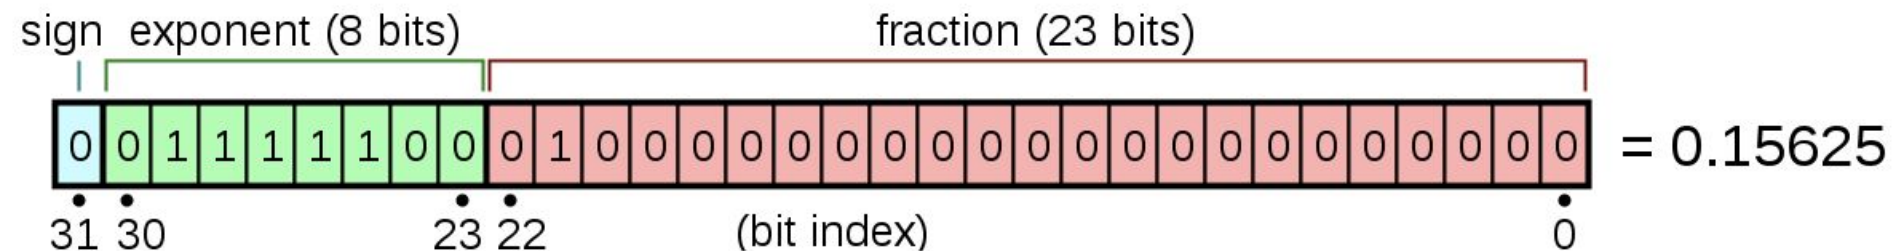
tuple of ints

int

np.array of ints

Floating point numbers

- Computer math is almost always floating point
- Like scientific notation with powers of 2 only



- np.float32 holds ~7 decimal digits
- np.float64 holds ~16 decimal digits
- Not every real number can be represented
- Too big = overflow, too small = underflow
- Only binary fractions (no exact 1/3, 1/5, etc)

```
print('Using f-strings, print 20 digits for each number')
for float_type in [np.float64, np.float32, np.float16]:
    x1 = float_type(1)
    print(f'{float_type} 1.0 is really {x1:.20f}')
    x2 = float_type(0.1)
    print(f'{float_type} 0.1 is really {x2:.20f}')
```

How does this code work?

What do you expect it to print?

```
print('Using f-strings, print 20 digits for each number')
for float_type in [np.float64, np.float32, np.float16]:
    x1 = float_type(1)
    print(f'{float_type} 1.0 is really {x1:.20f}')
    x2 = float_type(0.1)
    print(f'{float_type} 0.1 is really {x2:.20f}')
```

Using f-strings, print 20 digits for each number

float64 1.0 is really 1.000000000000000000000000

float64 0.1 is really 0.100000000000000000000555

float32 1.0 is really 1.000000000000000000000000

float32 0.1 is really 0.10000000149011611938

float16 1.0 is really 1.000000000000000000000000

float16 0.1 is really 0.09997558593750000000

Floating-point operations (flops)

- A computer can do billions/second (Gflops)
 - This is why we tolerate floating point issues
- Floating point arithmetic has round-off error
- Swamping ($A+B \approx A$) or Cancelation ($A-B \approx 0$) are the most common types of round-off error

What values of A, B might cause swamping?

What values might cause cancelation?

Try it in [Google Colab](#)!

Noticing round-off error

1. Avoid subtracting very similar numbers

Why?

2. Avoid adding big + small or dividing big/small
3. Avoid multiplying big*big or small*small
4. Check your answer if possible
5. Test for “nearness” instead of exact equality

Noticing round-off error

1. Avoid subtracting very similar numbers
 - $\text{sqrt}(x + 1) - \text{sqrt}(x) = 1 / (\text{sqrt}(x + 1) + \text{sqrt}(x))$
2. Avoid adding big + small or dividing big/small
 - $\text{big} / \text{small} \approx \text{big} / (\text{small} + \text{epsilon})$
3. Avoid multiplying big*big or small*small
 - $\log(A*A) = 2 \log(A)$
4. Test your answer if possible
 - if solving for $f(x) = 0$, check $f(x)$
5. Test for close instead of exact equality
 - $\text{abs}(f(x)) < 0.001$, not $f(x) == 0.0$

Dealing with round-off error

Try an expression analyzer such as

<https://herbie.uwplse.org/demo/>

Improves accuracy by testing alternate expressions over a random sample of inputs



Herbie web demo

Write a formula below, and Herbie will try to improve it. Enter approximate ranges for inputs.

[Show an example](#) | [Use FPCore](#)

$\sqrt{x + 1} - \sqrt{x}$

x: 0 to 1.79e308

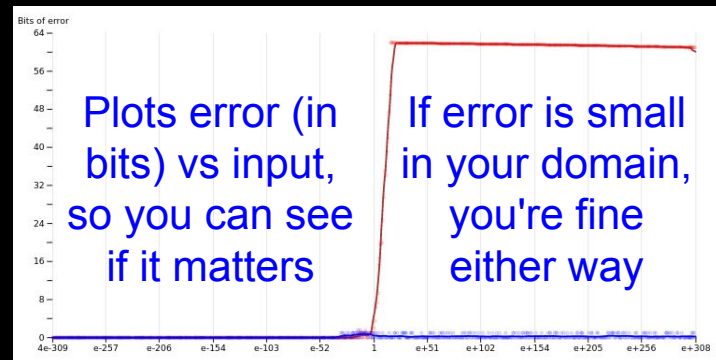
Improve with Herbie

$$\sqrt{x + 1} - \sqrt{x}$$

↓

$$\frac{1}{\sqrt{1 + x} + \sqrt{x}}$$

Can output math
or code, including
Python



Which float do you get by default?

```
x = 0.1 # default Python float
print(f'{type(x).__name__} 0.1 is really {x:.20f}')
xx = np.array([0.1]) # default Numpy array
print(f'{type(xx[0]).__name__} 0.1 is really {xx[0]:.20f}')
```

How does this code work?

What do you expect it to print?

Which float do you get by default?

```
x = 0.1 # default Python float
print(f'{type(x).__name__} 0.1 is really {x:.20f}')
xx = np.array([0.1]) # default Numpy array
print(f'{type(xx[0]).__name__} 0.1 is really {xx[0]:.20f}')
```

float 0.1 is really 0.10000000000000000000555

float64 0.1 is really 0.10000000000000000000555

Python and Numpy both use float64 by default
(most precise float type implemented in hardware,
thus most precision available while keeping speed)

NumPy array examples

Try the following exercises from PNM 2.8, 17-21:

17. Generate an array with size 100 evenly spaced between -10 to 10 using *linspace* function in Numpy.
18. Let array_a be an array [-1, 0, 1, 2, 0, 3]. Write a command that will return an array consisting of all the elements of array_a that are larger than zero.
Hint: Use logical expression as the index of the array.
19. Create an array $y = \begin{pmatrix} 3 & 5 & 3 \\ 2 & 2 & 5 \\ 3 & 8 & 9 \end{pmatrix}$ and calculate the transpose of the array.
20. Create a zero array with size (2, 4).
21. Change the 2nd column in the above array to 1.

Next reading: Taylor's Theorem (PNM 18.3),
Real Analysis (F&B 1.2),